

# CS 267: Introduction to Parallel Machines and Programming Models Lecture 3

James Demmel

[www.cs.berkeley.edu/~demmel/cs267\\_Spr14/](http://www.cs.berkeley.edu/~demmel/cs267_Spr14/)

CS267 Lecture 3

1

## Outline


- Overview of parallel machines (~hardware) and programming models (~software)
  - Shared memory
  - Shared address space
  - Message passing
  - Data parallel
  - Clusters of SMPs or GPUs
  - Grid
- Note: Parallel machine may or may not be tightly coupled to programming model
  - Historically, tight coupling
  - Today, portability is important

01/28/2014

CS267 Lecture 3

2

## A generic parallel architecture



- Where is the memory physically located?
- Is it connected directly to processors?
- What is the connectivity of the network?

01/28/2014

CS267 Lecture 3

3

## Parallel Programming Models

- **Programming model** is made up of the languages and libraries that create an abstract view of the machine
- Control
  - How is parallelism **created**?
  - What **orderings** exist between operations?
- Data
  - What data is **private** vs. **shared**?
  - How is logically shared data accessed or **communicated**?
- Synchronization
  - What operations can be used to coordinate parallelism?
  - What are the **atomic** (indivisible) operations?
- Cost
  - How do we account for the **cost** of each of the above?

01/28/2014

CS267 Lecture 3

4

## Simple Example


- Consider applying a function **f** to the elements of an array **A** and then computing its sum:

$$\sum_{i=0}^{n-1} f(A[i])$$

- Questions:

- Where does **A** live? All in single memory? Partitioned?
- What work will be done by each processors?
- They need to coordinate to get a single result, how?

**A** = array of all data  
**fA** = f(**A**)  
**s** = sum(**fA**)




01/28/2014

CS267 Lecture 3

5

## Programming Model 1: Shared Memory

- Program is a collection of threads of control.
  - Can be created dynamically, mid-execution, in some languages
- Each thread has a set of **private variables**, e.g., local stack variables
- Also a set of **shared variables**, e.g., static variables, shared common blocks, or global heap.
  - Threads communicate **implicitly** by writing and reading shared variables.
  - Threads coordinate by **synchronizing** on shared variables



01/28/2014

CS267 Lecture 3

6


## Simple Example

- Shared memory strategy:
  - small number  $p \ll n = \text{size}(A)$  processors
  - attached to single memory
- Parallel Decomposition:
  - Each evaluation and each partial sum is a task.
- Assign  $n/p$  numbers to each of  $p$  procs
  - Each computes independent "private" results and partial sum.
  - Collect the  $p$  partial sums and compute a global sum.

$$\sum_{i=0}^{n-1} f(A[i])$$

Two Classes of Data:

- Logically Shared
  - The original  $n$  numbers, the global sum.
- Logically Private
  - The individual function evaluations.
  - What about the individual partial sums?



01/28/2014

CS267 Lecture 3

7

## Shared Memory "Code" for Computing a Sum

```
fork(sum, a[0:n/2-1]);
sum(a[n/2, n-1]);
```

```
static int s = 0;
```

Thread 1

```
for i = 0, n/2-1
  s = s + f(A[i])
```

Thread 2

```
for i = n/2, n-1
  s = s + f(A[i])
```

- What is the problem with this program?
- A **race condition** or **data race** occurs when:
  - Two processors (or two threads) access the same variable, and at least one does a write.
  - The accesses are concurrent (not synchronized) so they could happen simultaneously

01/28/2014

CS267 Lecture 3

8

## Shared Memory “Code” for Computing a Sum

A = 

3	5
---	---

 $f(x) = x^2$

static int s = 0;

Thread 1		Thread 2	
....		...	
compute f([A[i]]) and put in reg0	9	compute f([A[i]]) and put in reg0	25
reg1 = s	0	reg1 = s	0
reg1 = reg1 + reg0	9	reg1 = reg1 + reg0	25
s = reg1	9	s = reg1	25
...		...	

- Assume A = [3,5],  $f(x) = x^2$ , and s=0 initially
- For this program to work, s should be  $3^2 + 5^2 = 34$  at the end
  - but it may be 34, 9, or 25
- The *atomic* operations are reads and writes
  - Never see ½ of one number, but += operation is *not* atomic
  - All computations happen in (private) registers

01/28/2014

CS267 Lecture 3

9

## Improved Code for Computing a Sum

static int s = 0;  
static lock lk;

Why not do lock  
inside loop?

Thread 1	Thread 2
local_s1 = 0 for i = 0, n/2-1 local_s1 = local_s1 + f(A[i]) lock(lk); s = s + local_s1 unlock(lk);	local_s2 = 0 for i = n/2, n-1 local_s2 = local_s2 + f(A[i]) lock(lk); s = s + local_s2 unlock(lk);

- Since addition is associative, it's OK to rearrange order
- Most computation is on private variables
  - Sharing frequency is also reduced, which might improve speed
  - But there is still a race condition on the update of shared s
  - The race condition can be fixed by adding locks (only one thread can hold a lock at a time; others wait for it)

01/28/2014

CS267 Lecture 3

10

## Review so far and plan for Lecture 3

### Programming Models

#### 1. Shared Memory

#### 2. Message Passing

#### 2a. Global Address Space

#### 3. Data Parallel

#### 4. Hybrid

### Machine Models

#### 1a. Shared Memory

#### 1b. Multithreaded Procs.

#### 1c. Distributed Shared Mem.

#### 2a. Distributed Memory

#### 2b. Internet & Grid Computing

#### 2c. Global Address Space

#### 3a. SIMD

#### 3b. Vector

#### 4. Hybrid


01/28/2014

CS267 Lecture 3

11

## Machine Model 1a: Shared Memory

- Processors all connected to a large shared memory.
  - Typically called **Symmetric Multiprocessors (SMPs)**
  - SGI, Sun, HP, Intel, IBM SMPs
  - Multicore chips, except that all caches are shared
- Advantage: uniform memory access (UMA)
- Cost: much cheaper to access data in cache than main memory
- Difficulty scaling to large numbers of processors
  - $\leq 32$  processors typical



Note: \$ = cache

01/28/2014

CS267 Lecture 3

12

## Problems Scaling Shared Memory Hardware

- Why not put more processors on (with larger memory?)
  - The memory bus becomes a bottleneck
  - Caches need to be kept *coherent*
- Example from a Parallel Spectral Transform Shallow Water Model (PSTSWM) demonstrates the problem
  - Experimental results (and slide) from Pat Worley at ORNL
  - This is an important kernel in atmospheric models
    - 99% of the floating point operations are multiplies or adds, which generally run well on all processors
    - But it does sweeps through memory with little reuse of operands, so uses bus and shared memory frequently
  - These experiments show performance per processor, with one “copy” of the code running independently on varying numbers of procs
    - The best case for shared memory: no sharing
    - But the data doesn’t all fit in the registers/cache

01/28/2014

CS267 Lecture 3


13

## Example: Problem in Scaling Shared Memory

### PSTSWM Sensitivity to Contention

- Performance degradation is a “smooth” function of the number of processes.
- No shared data between them, so there should be perfect parallelism.

- (Code was run for a 18 vertical levels with a range of horizontal sizes.)



Process scaling on IBM p690

OAK RIDGE NATIONAL LABORATORY  
U. S. DEPARTMENT OF ENERGY

UT-BATTELLE


01/28/2014

CS267 Lecture 3

From Pat Worley, ORNL 14

## Machine Model 1b: Multithreaded Processor

- Multiple thread “contexts” without full processors
- Memory and some other state is shared
- Sun Niagara processor (for servers)
  - Up to 64 threads all running simultaneously (8 threads x 8 cores)
  - In addition to sharing memory, they share floating point units
  - Why? Switch between threads for long-latency memory operations
- Cray MTA and Eldorado processors (for HPC)

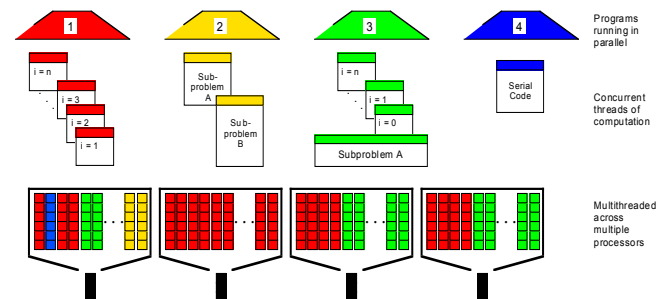


01/28/2014

CS267 Lecture 3

15

## Eldorado Processor (logical view)




01/28/2014

CS267 Lecture 3

Source: John Feo, Cray 16

### Machine Model 1c: Distributed Shared Memory

- Memory is logically shared, but physically distributed
  - Any processor can access any address in memory
  - Cache lines (or pages) are passed around machine
- SGI is canonical example (+ research machines)
  - Scales to 512 (SGI Altix (Columbia) at NASA/Ames)
  - Limitation is *cache coherency protocols* – how to keep cached copies of the same address consistent



01/28/2014

CS267 Lecture 3

17

### Review so far and plan for Lecture 3

Programming Models	Machine Models
1. Shared Memory	1a. Shared Memory 1b. Multithreaded Procs. 1c. Distributed Shared Mem.
2. Message Passing	2a. Distributed Memory 2b. Internet & Grid Computing
2a. Global Address Space	2c. Global Address Space
3. Data Parallel	3a. SIMD 3b. Vector
4. Hybrid	4. Hybrid

01/28/2014

CS267 Lecture 3

18

### Review so far and plan for Lecture 3

Programming Models	Machine Models
1. Shared Memory	1a. Shared Memory 1b. Multithreaded Procs. 1c. Distributed Shared Mem.
2. Message Passing	2a. Distributed Memory 2b. Internet & Grid Computing
2a. Global Address Space	2c. Global Address Space
3. Data Parallel	3a. SIMD 3b. Vector
4. Hybrid	4. Hybrid


01/28/2014

CS267 Lecture 3

19

### Programming Model 2: Message Passing

- Program consists of a collection of **named** processes.
  - Usually fixed at program startup time
  - Thread of control plus local address space -- **NO shared data.**
  - Logically shared data is partitioned over local processes.
- Processes communicate by explicit send/receive pairs
  - Coordination is implicit in every communication event.
  - MPI (Message Passing Interface) is the most commonly used SW



01/28/2014

CS267 Lecture 3

20

## Computing $s = f(A[1]) + f(A[2])$ on each processor

° First possible solution – what could go wrong?

Processor 1  
 $x_{\text{local}} = f(A[1])$   
 send  $x_{\text{local}}$ , proc2  
 receive  $x_{\text{remote}}$ , proc2  
 $s = x_{\text{local}} + x_{\text{remote}}$

Processor 2  
 $x_{\text{local}} = f(A[2])$   
 send  $x_{\text{local}}$ , proc1  
 receive  $x_{\text{remote}}$ , proc1  
 $s = x_{\text{local}} + x_{\text{remote}}$

° If send/receive acts like the telephone system? The post office?

° Second possible solution

Processor 1  
 $x_{\text{local}} = f(A[1])$   
 send  $x_{\text{local}}$ , proc2  
 receive  $x_{\text{remote}}$ , proc2  
 $s = x_{\text{local}} + x_{\text{remote}}$

Processor 2  
 $x_{\text{local}} = f(A[2])$   
 receive  $x_{\text{remote}}$ , proc1  
 send  $x_{\text{local}}$ , proc1  
 $s = x_{\text{local}} + x_{\text{remote}}$

° What if there are more than 2 processors?

01/28/2014

CS267 Lecture 3

21

## MPI – the de facto standard

MPI has become the de facto standard for parallel computing using message passing

### Pros and Cons of standards

- MPI created finally a standard for applications development in the HPC community → portability
- The MPI standard is a least common denominator building on mid-80s technology, so may discourage innovation

Programming Model reflects hardware!

*“I am not sure how I will program a Petaflops computer, but I am sure that I will need MPI somewhere” – HDS 2001*


01/28/2014

CS267 Lecture 3

22

## Machine Model 2a: Distributed Memory

- Cray XE6 (Hopper)
- PC Clusters (Berkeley NOW, Beowulf)
- Edison, Hopper, Franklin, IBM SP-3, Millennium, are distributed memory machines, but the **nodes** are SMPs.
- Each processor has its own memory and cache but cannot directly access another processor's memory.
- Each “node” has a Network Interface (NI) for all communication and synchronization.



01/28/2014

CS267 Lecture 3

23

## PC Clusters: Contributions of Beowulf

- An experiment in parallel computing systems (1994)
- Established vision of low cost, high end computing
  - Cost effective because it uses off-the-shelf parts
- Demonstrated effectiveness of PC clusters for some (not all) classes of applications
- Provided networking software
- Conveyed findings to broad community (great PR)
- Tutorials and book
- Design standard to rally community!
- Standards beget: books, trained people, software ... virtuous cycle



Adapted from Gordon Bell, presentation at Salishan

01/28/2014

CS267 Lecture 3

24

## Tflop/s and Pflop/s Clusters (2009 data)

The following are examples of clusters configured out of separate networks and processor components

- About 82% of Top 500 are clusters (Nov 2009, up from 72% in 2005),
  - 4 of top 10
- IBM Cell cluster at Los Alamos (Roadrunner) is #2
  - 12,960 Cell chips + 6,948 dual-core AMD Opterons;
    - 129600 cores altogether
  - 1.45 PFlops peak, 1.1PFlops Linpack, 2.5MWatts
  - Infiniband connection network
- For more details use "database/sublist generator" at [www.top500.org](http://www.top500.org)

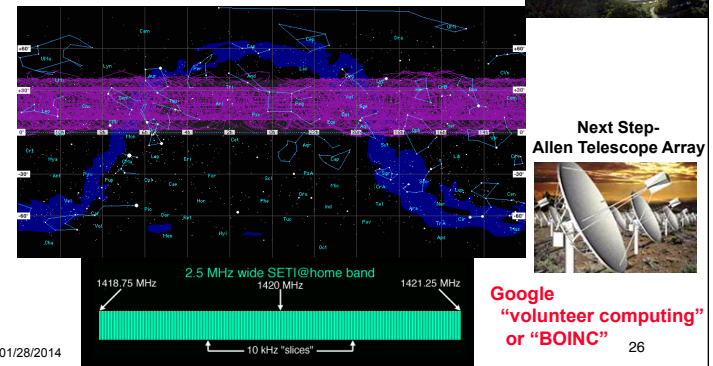
01/28/2014

CS267 Lecture 3

25

## Machine Model 2b: Internet/Grid Computing

- **SETI@Home**: Running on 3.3M hosts, 1.3M users (1/2013)
  - ~1000 CPU Years per Day (older data)
  - 485,821 CPU Years so far
- Sophisticated Data & Signal Processing Analysis
- Distributes Datasets from Arecibo Radio Telescope




01/28/2014

26

## Programming Model 2a: Global Address Space

- Program consists of a collection of **named** threads.
  - Usually fixed at program startup time
  - Local and shared data, as in shared memory model
  - But, shared data is partitioned over local processes
  - Cost models says remote data is expensive
- Examples: UPC, Titanium, Co-Array Fortran
- Global Address Space programming is an intermediate point between message passing and shared memory




01/28/2014

CS267 Lecture 3

27

## Machine Model 2c: Global Address Space

- Cray T3D, T3E, X1, and HP Alphaser server cluster
- Clusters built with Quadrics, Myrinet, or Infiniband
- The network interface supports RDMA (Remote Direct Memory Access)
  - NI can directly access memory without interrupting the CPU
  - One processor can read/write memory with one-sided operations (put/get)
  - Not just a load/store as on a shared memory machine
    - Continue computing while waiting for memory op to finish
  - Remote data is typically not cached locally



01/28/2014

CS267 Lecture 3

28

## Review so far and plan for Lecture 3

Programming Models	Machine Models
1. Shared Memory	1a. Shared Memory 1b. Multithreaded Procs. 1c. Distributed Shared Mem.
2. Message Passing	2a. Distributed Memory 2b. Internet & Grid Computing
2a. Global Address Space	2c. Global Address Space
3. Data Parallel	3a. SIMD 3b. Vector
4. Hybrid	4. Hybrid


01/28/2014

CS267 Lecture 3

29

## Programming Model 3: Data Parallel

- Single thread of control consisting of **parallel operations**.
  - $A = B + C$  could mean add two arrays in parallel
- Parallel operations applied to all (or a defined subset) of a data structure, usually an array
  - Communication is implicit in parallel operators**
  - Elegant and easy to understand and reason about**
  - Coordination is implicit – statements executed synchronously**
  - Similar to Matlab language for array operations**
- Drawbacks:
  - Not all problems fit this model**
  - Difficult to map onto coarse-grained machines**




01/28/2014

CS267 Lecture 3

30

## Machine Model 3a: SIMD System

- A large number of (usually) small processors.
  - A single “control processor” issues each instruction.**
  - Each processor executes the same instruction.**
  - Some processors may be turned off on some instructions.**
- Originally machines were specialized to scientific computing, few made (CM2, Maspar)
- Programming model can be implemented in the compiler
  - mapping  $n$ -fold parallelism to  $p$  processors,  $n \gg p$ , but it's hard (e.g., HPF)**



01/28/2014

CS267 Lecture 3

31

## Machine Model 3b: Vector Machines

- Vector architectures are based on a single processor
  - Multiple functional units**
  - All performing the same operation**
  - Instructions may specify large amounts of parallelism (e.g., 64-way) but hardware executes only a subset in parallel**
- Historically important
  - Overtaken by MPPs in the 90s**
- Re-emerging in recent years
  - At a large scale in the Earth Simulator (NEC SX6) and Cray X1**
  - At a small scale in SIMD media extensions to microprocessors**
    - SSE, SSE2 (Intel: Pentium/IA64)
    - AltiVec (IBM/Motorola/Apple: PowerPC)
    - VIS (Sun: Sparc)
  - At a larger scale in GPUs**
- Key idea: Compiler does some of the difficult work of finding parallelism, so the hardware doesn't have to

01/28/2014


CS267 Lecture 3

32




## Vector Processors

- Vector instructions operate on a vector of elements
  - These are specified as operations on vector registers



- A supercomputer vector register holds ~32-64 elts
  - The number of elements is larger than the amount of parallel hardware, called vector pipes or lanes, say 2-4
- The hardware performs a full vector operation in
  - #elements-per-vector-register / #pipes



01/28/2014


CS267 Lecture 3

33

## Cray X1: Parallel Vector Architecture

Cray combines several technologies in the X1

- 12.8 Gflop/s Vector processors (MSP)
- Shared caches (unusual on earlier vector machines)
- 4 processor nodes sharing up to 64 GB of memory
- Single System Image to 4096 Processors
- Remote put/get between nodes (faster than MPI)

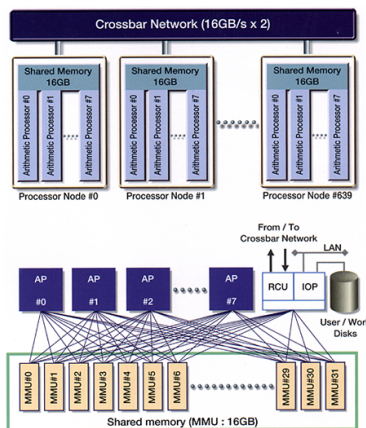


01/28/2014

CS267 Lecture 3

34

## Earth Simulator Architecture



### Parallel Vector Architecture

- High speed (vector) processors
- High memory bandwidth (vector architecture)
- Fast network (new crossbar switch)

Rearranging commodity parts can't match this performance

01/28/2014

CS267 Lecture 3

35

## Review so far and plan for Lecture 3

### Programming Models

- Shared Memory
- Message Passing
- Data Parallel
- Hybrid

### Machine Models

- Shared Memory
- Multithreaded Procs.
- Distributed Shared Mem.
- Distributed Memory
- Internet & Grid Computing
- Global Address Space
- SIMD & GPU
- Vector
- Hybrid

01/28/2014

CS267 Lecture 3

36

### Machine Model 4: Hybrid machines

- Multicore/SMPs are a building block for a larger machine with a network
- Old name:
  - CLUMP = Cluster of SMPs
- Many modern machines look like this:
  - Edison and Hopper (2x12 way nodes), most of Top500
- What is an appropriate programming model #4 ???
  - Treat machine as “flat”, always use message passing, even within SMP (simple, but ignores an important part of memory hierarchy).
  - Shared memory within one SMP, but message passing outside of an SMP.
- GPUs may also be building block
  - Nov 2013 Top500: 10% have accelerators, but 35% of performance

01/28/2014

CS267 Lecture 3

37

### Programming Model 4: Hybrids

- Programming models can be mixed
  - Message passing (MPI) at the top level with shared memory within a node is common
  - New DARPA HPCS languages mix data parallel and threads in a global address space
  - Global address space models can (often) call message passing libraries or vice versa
  - Global address space models can be used in a hybrid mode
    - Shared memory when it exists in hardware
    - Communication (done by the runtime system) otherwise
- For better or worse
  - Supercomputers often programmed this way for peak performance

01/28/2014

CS267 Lecture 3

38

### Review so far and plan for Lecture 3

#### Programming Models

1. Shared Memory

2. Message Passing

2a. Global Address Space

3. Data Parallel

4. Hybrid

#### Machine Models

1a. Shared Memory

1b. Multithreaded Procs.

1c. Distributed Shared Mem.

2a. Distributed Memory

2b. Internet & Grid Computing

2c. Global Address Space

3a. SIMD & GPU

3b. Vector

4. Hybrid

**What about GPU? What about Cloud?**

01/28/2014

CS267 Lecture 3

39

### What about GPU and Cloud?

- GPU's big performance opportunity is data parallelism
  - Most programs have a mixture of highly parallel operations, and some not so parallel
  - GPUs provide a threaded programming model (CUDA) for data parallelism to accommodate both
  - Current research attempting to generalize programming model to other architectures, for portability (OpenCL)
  - Guest lecture later in the semester
- Cloud computing lets large numbers of people easily share  $O(10^5)$  machines
  - MapReduce was first programming model: data parallel on distributed memory
  - More flexible models (Hadoop...) invented since then
  - Guest lecture later in the semester
- Both may be used for class projects

01/28/2014

CS267 Lecture 3

40

### Lessons from Lecture 3

- Three basic conceptual models
  - Shared memory
  - Distributed memory
  - Data paralleland hybrids of these machines
- All of these machines rely on dividing up work into parts that are:
  - Mostly independent (little synchronization)
  - About same size (load balanced)
  - Have good locality (little communication)
- Next Lecture: How to identify parallelism and locality in applications

01/28/2014

CS267 Lecture 3

41